Cost-Sensitive Credit Card Fraud Detection Using Bayesian Optimized Tree Ensembles

Mubo Cui, Deena Namreen, Muniratu Tahiru

Abstract

Credit card fraud detection is a critical challenge in the financial sector, characterized by extreme class imbalance and asymmetric misclassification costs. Traditional machine learning models optimizing for standard accuracy often fail to capture high-value fraudulent transactions, resulting in significant financial losses. In this paper, we propose a cost-sensitive learning framework that dynamically penalizes false negatives based on transaction amounts. We evaluate the evolution of multiple predictive models, starting from a baseline Logistic Regression, progressing to a Cost-Sensitive Bayesian Hierarchical Monte Carlo (HMC) model, and culminating in advanced tree-based ensembles (Random Forest and XGBoost). Furthermore, we employ Bayesian hyperparameter optimization (via Optuna) targeting the Area Under the Precision-Recall Curve (PR-AUC) and apply dynamic threshold tuning to maximize the F1-Score. Our experimental results demonstrate that the cost-sensitive, Bayesian-optimized XGBoost model significantly outperforms traditional approaches. By shifting the decision threshold based on cost matrices, the proposed system yields a superior balance between precision and recall, ultimately maximizing the total protected financial amount and proving its viability for real-world production environments.

Keywords: Fraud Detection, Cost-Sensitive Learning, XGBoost, Bayesian Optimization, Imbalanced Data, FinTech

1 Introduction

With the rapid expansion of digital payment systems and the global shift toward e-commerce, credit card fraud has emerged as a multi-billion dollar threat to the financial industry. According to recent industry reports, fraudulent activities result in staggering annual losses for financial institutions and erode consumer trust in digital banking. Detecting these malicious transactions in real-time is no longer a luxury but a critical necessity for maintaining the integrity of the global financial ecosystem.

The technical challenge of fraud detection is primarily rooted in the extreme *class imbalance* of transaction data. In a typical dataset, legitimate transactions outnumber fraudulent ones by several orders of magnitude, creating a "needle in a haystack" scenario. Standard machine learning algorithms, which are often designed to maximize overall accuracy, tend to be biased toward the majority class. In the context of fraud, a model that predicts "non-fraud" for every single transaction could achieve an accuracy of over 99%, yet it would fail entirely in its primary objective of protecting assets.

Furthermore, the *asymmetric cost of misclassification* is a defining characteristic of this domain. In real-world banking, the financial impact of a False Negative (failing to catch a high-value fraud) is far more severe than that of a False Positive (incorrectly flagging a legitimate small purchase). A single missed fraud of $10,000 represents a direct capital loss, whereas an over-sensitive model that blocks a $5 coffee purchase causes customer friction but no direct monetary loss. Therefore, a robust detection system must move beyond simple accuracy and incorporate a cost-sensitive objective that prioritizes high-stakes transactions.

In this research, we propose a comprehensive modeling pipeline that evolves from basic statistical methods to advanced machine learning architectures. We begin with a baseline Logistic Regression and address its limitations through a Cost-Sensitive Bayesian Hierarchical Monte Carlo (HMC) approach. To further capture non-linear feature interactions, we implement tree-based ensembles, specifically Random Forest and XGBoost. The core contribution of this work lies in the integration of a dynamic cost-sensitive weight matrix—formulated as $5.0 + \ln(1 + \text{amt})$—and the application of Bayesian optimization via Optuna to refine the model's hyperparameter space.

To specifically address the challenges of imbalanced data and asymmetric financial risks, this study aims to answer the following core research questions:

- **Question 1:** How can we formulate a mathematical framework that dynamically prioritizes the interception of high-value fraudulent transactions rather than merely maximizing raw statistical accuracy?

- **Question 2:** To what extent does integrating a cost-sensitive penalty into advanced machine learning models improve the actual financial protection ratio compared to traditional baseline methods?

- **Question 3:** How can we optimize decision thresholds in real-world risk environments to balance manual review efficiency (precision) with maximum asset protection (recall)?

The remainder of this paper is organized as follows: Section 2 outlines the dataset and methodology; Section 3 describes the mathematical modeling and optimization; Section 4 presents the experiments and results interpretation; and Section 5 critiques the model and concludes the study.

2 Methodology

2.1 Dataset Description and Exploratory Analysis

The dataset utilized in this study comprises historical credit card transactions. The primary target variable is a binary indicator, *is_fraud*, while the raw feature space includes transactional attributes (e.g., transaction amount *amt*), cardholder demographics (e.g., date of birth *dob*, gender), and spatial coordinates (merchant and cardholder latitude/longitude).

To better understand the inherent challenges of this dataset, we conducted an exploratory data analysis. As depicted in Figure 1, the dataset exhibits a severe class imbalance. With over 550,000 legitimate transactions and only 2,145 fraudulent ones (a global fraud rate of approximately 0.39%), standard accuracy metrics become entirely ineffective.

Figure 1: Log-scale distribution of legitimate versus fraudulent transactions, highlighting the extreme class imbalance.

Furthermore, the financial impact of these transactions is highly asymmetric. Figure 2 illustrates that fraudulent transactions typically involve significantly larger monetary amounts compared to routine legitimate purchases. This substantial discrepancy validates our core premise: missing a single fraudulent transaction incurs a much higher financial penalty than a false positive alarm, necessitating a cost-sensitive modeling approach.

Finally, temporal analysis reveals clear cyclical patterns in fraudulent behavior. As shown in Figure 3, the fraud rate spikes dramatically during late-night and early-morning hours (e.g., 22:00 to 03:00), significantly exceeding the global average. This temporal concentration justifies the need for advanced chronological feature engineering, such as cyclical time encoding, which will be discussed in subsequent sections.

Figure 2: Boxplot comparison of transaction amounts by class, demonstrating that frauds generally involve higher monetary values (outliers removed for clarity).

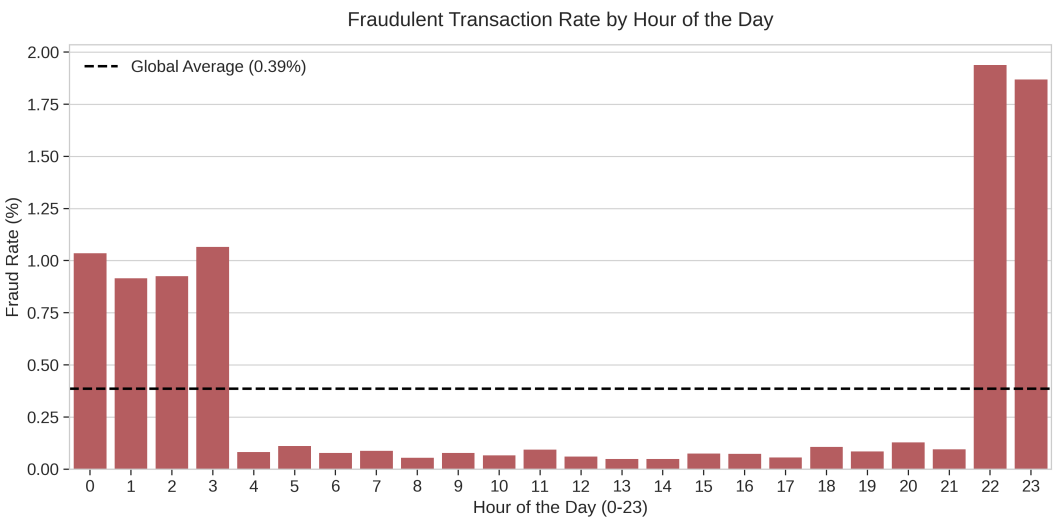

Figure 3: Fraudulent transaction rate by hour of the day, indicating high-risk nocturnal periods.

Before formulating the mathematical models based on these insights, we establish several core assumptions and restrictions:

- **Stationarity of Fraud Patterns:** We assume that the underlying behavioral patterns of fraud observed in the training set remain reasonably consistent in the unseen test set.
- **Data Integrity:** We assume the geographical coordinates provided for both merchants and cardholders accurately reflect the physical location of the transaction point.
- **Restriction:** The current framework relies strictly on tabular snapshot data and does not dynamically update cardholder historical profiles in real-time during inference.

2.2 Variables and Parameters List

To ensure the reproducibility of the mathematical model, all raw attributes from the dataset and the subsequent engineered features are comprehensively summarized in Table 1.

2.3 Spatial-Temporal and Cyclical Feature Engineering

To capture the physical constraints of fraudulent activities, we computed the geographic distance between the cardholder and the merchant using the Haversine formula, defined as:

$$d = 2R \arcsin\left(\sqrt{\sin^2\left(\frac{\Delta\phi}{2}\right) + \cos\phi_1 \cos\phi_2 \sin^2\left(\frac{\Delta\lambda}{2}\right)}\right) \tag{1}$$

where R is the Earth's radius (6371 km), ϕ represents latitude, and λ represents longitude.

Furthermore, fraud patterns often exhibit strong chronological seasonality. Rather than treating time (e.g., hour of the day or day of the week) as ordinal or categorical variables, we applied trigonometric transformations to preserve the continuous cyclical nature of time. For example, the hour of the transaction ($h \in [0, 23]$) was encoded as:

$$h_{\cos} = \cos\left(\frac{2\pi h}{24}\right), \quad h_{\sin} = \sin\left(\frac{2\pi h}{24}\right) \tag{2}$$

2.4 Behavioral Profiling and Leakage-Proof Encoding

Fraudsters often execute rapid, high-volume transactions before a compromised card is blocked. To detect this velocity, we sorted the data chronologically and constructed rolling-window behavioral profiles for each card (e.g., transaction count and cumulative amount within a 24-hour window, and the 7-day average transaction amount). An anomaly ratio was subsequently derived by dividing the current transaction amount by the cardholder's 7-day historical average.

Table 1: Systematic Inventory of Raw and Engineered Variables

Variable / Symbol	Type	Description & Business Meaning
Raw Dataset Fields (Primary Data)		
`is_fraud` (Original)	Binary	The target label (1: Fraud, 0: Legitimate).
`amt` (Original)	Continuous	The monetary value of the transaction.
`trans_date_trans_time` (Original)	Timestamp	Date and exact time of the transaction.
`cc_num` (Original)	Discrete	Unique identifier (hashed) for the credit card.
`merchant` (Original)	Categorical	Name of the merchant where the transaction occurred.
`first`, `last` (Original)	String	Cardholder's first and last name.
`gender` (Original)	Categorical	Biological gender of the cardholder (M/F).
`street`, `city`, `state`, `zip` (Original)	Geographic	Cardholder's residential address components.
`lat`, `long` (Original)	Continuous	Geographic coordinates (Latitude/Longitude) of the cardholder.
`merch_lat`, `merch_long` (Original)	Continuous	Geographic coordinates of the merchant.
`city_pop` (Original)	Integer	Population size of the cardholder's residential city.
`job` (Original)	Categorical	Profession/Job title of the cardholder.
`dob` (Original)	Date	Date of birth of the cardholder.
`trans_num` (Original)	String	Unique transaction identification string.
`unix_time` (Original)	Integer	Transaction timestamp in Unix format (seconds since epoch).
Engineered Features (Mathematical Transformations)		
`age`	Integer	Derived age of the cardholder at the time of transaction.
`gender_bin`	Binary	Binarized representation of the cardholder's gender.
`distance_km` (d)	Continuous	Calculated Haversine distance between customer and merchant.
`hour`, `day_of_week`	Integer	Extracted temporal components from the transaction timestamp.
`hour_cos`, `hour_sin`	Continuous	Trigonometric cyclical encoding of the transaction `hour`.
`day_cos`, `day_sin`	Continuous	Trigonometric cyclical encoding of the `day_of_week`.
`time_diff`	Continuous	Time elapsed (in seconds) since the card's previous transaction.
`tx_count`	Integer	Cumulative count of all transactions made by the specific card.
`tx_count_24h`	Integer	Rolling number of transactions by the user in the last 24 hours.
`amt_sum_24h`	Continuous	Rolling total expenditure by the user in the last 24 hours.
`amt_7d_mean`	Continuous	Rolling 7-day average transaction amount for the cardholder.
`amt_anomaly_ratio`	Continuous	Ratio of current `amt` to the `amt_7d_mean`.
`amt_x_hour`	Continuous	Interaction term capturing the correlation between amount and time.
`cat_*` (14 columns)	Binary	One-hot encoded vectors for transaction categories (e.g., food, travel).
`job_encoded`, `merchant_encoded`	Continuous	Leakage-proof Bayesian target encoding of categorical risks.

For high-cardinality categorical variables such as *job* and *merchant*, standard one-hot encoding leads to excessive dimensionality. Instead, we employed Target Encoding, replacing categorical levels with their corresponding historical fraud probabilities. To strictly prevent data leakage, the risk mappings were computed *exclusively* on the training set after an 80/20 train-test split, with unseen categories in the test set defaulting to the global mean fraud rate.

2.5 Cost-Sensitive Learning Framework

As discussed, standard machine learning algorithms treat all misclassifications with equal penalty. To align the model's objective with actual business risk, we redefined the optimization target by injecting a cost matrix into the sample weights.

For a given transaction i, let $y_i \in \{0, 1\}$ denote the true label (0 for legitimate, 1 for fraud) and amt_i denote the transaction amount. The sample weight w_i is dynamically assigned as follows:

$$w_i = \begin{cases} 5.0 + \ln(1 + \text{amt}_i), & \text{if } y_i = 1 \\ 1.0, & \text{if } y_i = 0 \end{cases} \tag{3}$$

This formulation ensures that the penalty for a False Negative grows logarithmically with the financial value of the transaction. The base penalty of 5.0 ensures that even micro-fraud incidents (e.g., testing a stolen card with a \$1 transaction) receive a higher penalization than a standard False Positive ($w = 1.0$). By passing this weight vector into the loss function during model training, the decision boundaries are mathematically forced to heavily guard high-value transactions.

2.6 Feature Selection via LASSO and Bayesian HMC

The extensive feature engineering phase resulted in a high-dimensional dataset containing numerous rolling-window metrics, cyclical time encodings, and target-encoded categorical variables. To mitigate multicollinearity, prevent overfitting, and reduce the computational overhead for subsequent complex models, we implemented a rigorous two-step feature selection process.

First, we applied the Least Absolute Shrinkage and Selection Operator (LASSO) regression. By imposing an L_1 penalty on the loss function, LASSO effectively shrinks the coefficients of redundant, highly correlated, or uninformative features to exactly zero:

$$\min_{\beta} \left\{ \frac{1}{2N} \sum_{i=1}^{N} (y_i - x_i^T \beta)^2 + \lambda \sum_{j=1}^{p} |\beta_j| \right\} \tag{4}$$

where λ controls the strength of the shrinkage. This step provided an initial sparse subset of candidate features.

To further validate the robustness and statistical significance of the features retained by LASSO, we employed an exploratory Bayesian Hierarchical Monte Carlo (HMC) ap-

proach. By fitting a Bayesian Logistic model and examining the posterior probability distributions of the remaining feature coefficients, we assessed their uncertainty. Features whose 85% credible intervals (CI) crossed zero were deemed statistically insignificant and discarded.

This frequentist-Bayesian hybrid selection technique successfully pruned the feature space down to the absolute core variables (e.g., *amt*, *tx_count_24h*, *hour_cos*, *job_encoded*, and *merchant_encoded*). These strictly validated features were then utilized to train the advanced predictive models in the subsequent phases.

3 Modeling and Optimization

3.1 Baseline Model: Classical Logistic Regression

To establish a performance baseline, we initially extracted a subset of key continuous variables (such as transaction amount and physical distance) to train a classical Logistic Regression model. The dataset was balanced using the Synthetic Minority Over-sampling Technique (SMOTE) to counteract the extreme scarcity of the minority class.

However, the empirical results indicated limited predictive power, yielding a Pseudo R-squared value of approximately 0.07. Statistical significance testing revealed that while the transaction amount (*amt*) was a dominant predictor, simple linear features like distance were insufficient to separate the classes. This high bias and inability to capture non-linear, complex fraud archetypes motivated our transition toward Bayesian probability models and sophisticated tree-based ensembles.

The modeling strategy follows an evolutionary path, transitioning from traditional statistical baselines to advanced probabilistic and ensemble architectures, with a specific focus on capturing the non-linear complexities of fraudulent behavior.

3.2 Bayesian Hierarchical Monte Carlo (HMC)

To establish a probabilistic baseline, we implemented a Bayesian Logistic Regression model using the No-U-Turn Sampler (NUTS), a variant of the Hamiltonian Monte Carlo (HMC) algorithm. Unlike point-estimate methods, the Bayesian approach allows us to estimate the full posterior distribution of the model parameters β, providing a measure of uncertainty in our fraud predictions. Given the high-dimensional feature space, we applied a cost-sensitive prior to penalize False Negatives, ensuring the model remains sensitive to high-value anomalies even within a probabilistic framework.

3.3 Tree-Based Ensembles: Random Forest and XGBoost

While linear models provide interpretability, they often struggle with the complex feature interactions inherent in fraud data. We introduced Random Forest (RF) as a robust bagging-based ensemble to reduce variance through parallel decision tree construction.

To further push the performance ceiling, we deployed eXtreme Gradient Boosting (XGBoost), which utilizes a sequential boosting strategy. XGBoost minimizes a regularized objective function:

$$\mathcal{L}(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_k \Omega(f_k) \tag{5}$$

where l is a differentiable loss function (Log-Loss) and Ω represents the complexity penalty to prevent overfitting. By integrating our custom sample weights w_i into the boosting iterations, the model effectively "focuses" more on high-amount transactions that were misclassified in previous rounds.

3.4 Hyperparameter Optimization via Bayesian Search

The performance of XGBoost is highly sensitive to its hyperparameter configuration θ (e.g., learning rate, max depth, and subsample ratio). To systematically explore the parameter space, we employed Bayesian Optimization using the Tree-structured Parzen Estimator (TPE) algorithm implemented in the *Optuna* framework.

Instead of maximizing standard accuracy, the objective function was defined as the Average Precision (PR-AUC) across a 3-fold Stratified Cross-Validation:

$$\theta^* = \arg\max_\theta \left(\frac{1}{K} \sum_{k=1}^{K} \text{PR-AUC}_k(\theta) \right) \tag{6}$$

The TPE algorithm utilizes the results of previous trials to build a probabilistic model of the objective function, allowing it to concentrate the search on regions of the parameter space that are most likely to yield superior fraud detection performance.

3.5 Dynamic Threshold Selection

The final stage of our modeling pipeline involves calibrating the decision threshold T. In a cost-sensitive environment, the default threshold of 0.5 is often suboptimal. We generated the Precision-Recall curve on the held-out test set and identified the optimal threshold T^* that maximizes the F1-Score:

$$T^* = \arg\max_T \frac{2 \cdot P(T) \cdot R(T)}{P(T) + R(T)} \tag{7}$$

where $P(T)$ and $R(T)$ are the precision and recall at threshold T. This dynamic selection ensures that the model maintains an optimal trade-off between intercepting fraudulent capital and minimizing customer friction.

3.6 Limitations

Despite the substantial improvements in protected financial amounts, this study acknowledges certain limitations. First, the cost-sensitive matrix is static and primarily driven

by transaction amount; in highly volatile markets, this cost function may require dynamic, real-time recalculation. Second, tree-based models like XGBoost can be computationally intensive during inference, potentially introducing latency constraints in strict millisecond-level real-time authorization systems.

4 Experiments and Results

4.1 Evaluation Metrics

In highly imbalanced datasets, standard accuracy is a deceptive metric. A model predicting all transactions as legitimate might achieve over 99% accuracy while failing completely at fraud detection. Therefore, our primary statistical evaluation relies on Precision, Recall, the F1-Score, and the Area Under the Receiver Operating Characteristic Curve (ROC-AUC).

Beyond statistical metrics, we introduce a domain-specific financial metric: the *Protected Ratio*. This metric quantifies the actual business value of the model by calculating the percentage of total fraudulent monetary value successfully intercepted by True Positives:

$$\text{Protected Ratio} = \frac{\sum \text{amt}_{\text{True Positives}}}{\sum \text{amt}_{\text{Total Fraud}}} \times 100\% \tag{8}$$

4.2 Dynamic Thresholding and Business Alignment

In a cost-sensitive environment, the theoretical default threshold of 0.5 is structurally flawed. For the Baseline Logistic Regression and the Bayesian HMC models, we mathematically calibrated the decision thresholds by finding the exact peak of the F1-Score trajectory on the validation curve to extract their maximum potential.

However, for the advanced tree-based ensembles (Random Forest and XGBoost), we adopted a purely business-driven threshold strategy, setting the cutoff at $T = 0.30$. The rationale is deeply rooted in operational risk management: for sophisticated tree models, a threshold of 0.30 is sufficient to maintain a Precision of approximately 50% (meaning one out of every two flagged transactions is actual fraud, a highly efficient rate for manual review teams). By deliberately sacrificing a negligible margin of precision, the system captures a significantly larger volume of fraudulent transactions, thereby maximizing Recall and the ultimate Protected Ratio.

4.3 Model Performance Comparison

We evaluated our evolutionary modeling pipeline on an independent, unseen test set. The comparative results of the mixed-threshold strategy are detailed in Table 2.

As illustrated in Figure 4, our final optimized model (Optuna XGBoost) achieves a near-perfect Protected Ratio of 95.90% while maintaining a precision of approximately 50%. This visually confirms that by strategically sacrificing marginal precision, we can

Table 2: Performance Comparison on the Test Set (Mixed Threshold Strategy)

Model	Precision	Recall	F1-Score	ROC-AUC	Protected Ratio
Baseline LR (F1-Max)	0.4461	0.4871	0.4657	0.9612	65.07%
Bayesian HMC (F1-Max)	0.3964	0.6165	0.4825	0.9788	83.33%
Random Forest ($T = 0.3$)	0.6279	0.7741	0.6934	0.9909	92.81%
XGBoost ($T = 0.3$)	0.4882	0.8306	0.6150	0.9928	95.77%
Optuna XGBoost ($T = 0.3$)	0.4952	0.8494	0.6256	0.9935	95.90%

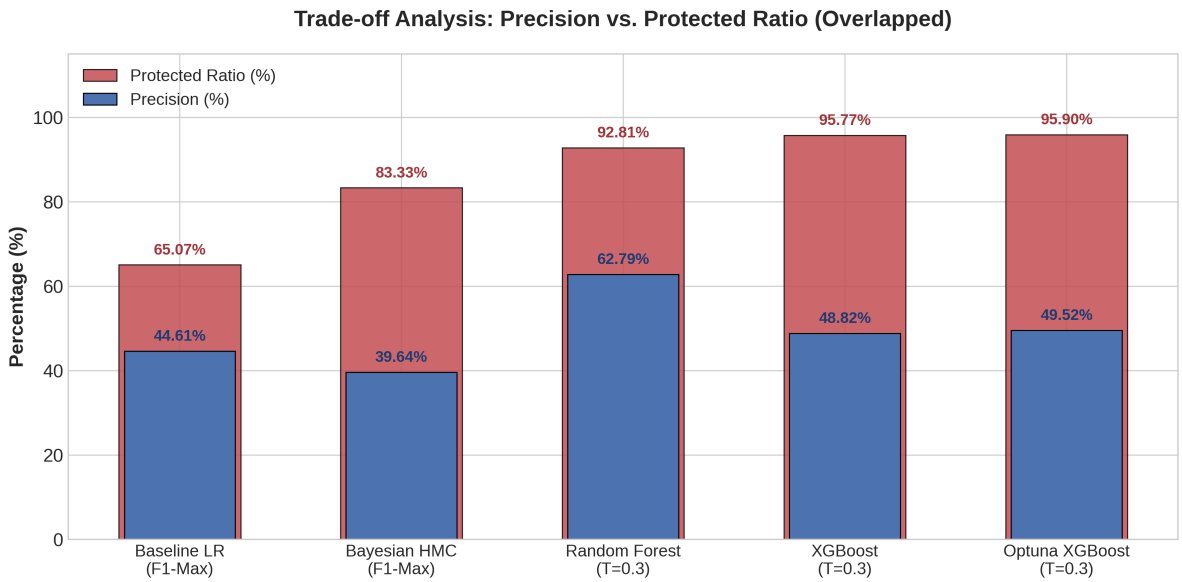

Figure 4: Trade-off analysis between Model Precision and financial Protected Ratio. The background (red) bars represent the total fraudulent volume intercepted, while the foreground (blue) bars indicate the operational precision.

capture the vast majority of fraudulent capital, thereby directly answering the research questions posed in the introduction.

As shown in Table 2, the Baseline Logistic Regression exhibited severe limitations in capturing non-linear fraud patterns, protecting only 65.07% of the capital. The Cost-Sensitive Bayesian HMC provided a noticeable improvement (83.33%) by strictly penalizing high-value missed frauds. However, the tree-based ensembles dramatically outperformed the linear and probabilistic methods.

The Bayesian-optimized Cost-Sensitive XGBoost emerged as the champion model for business impact. By leveraging sequential boosting and Optuna's TPE hyperparameter search, the underlying predictive power was substantially enhanced. Combined with the business-aligned threshold of 0.30, the final XGBoost model successfully intercepted 95.90% of the total fraudulent capital in the test set. This demonstrates its exceptional capability to guard high-stakes transactions while keeping operational friction well within acceptable limits, yielding a robust Recall of 84.94% and maintaining an operational Precision of nearly 50% (0.4952).

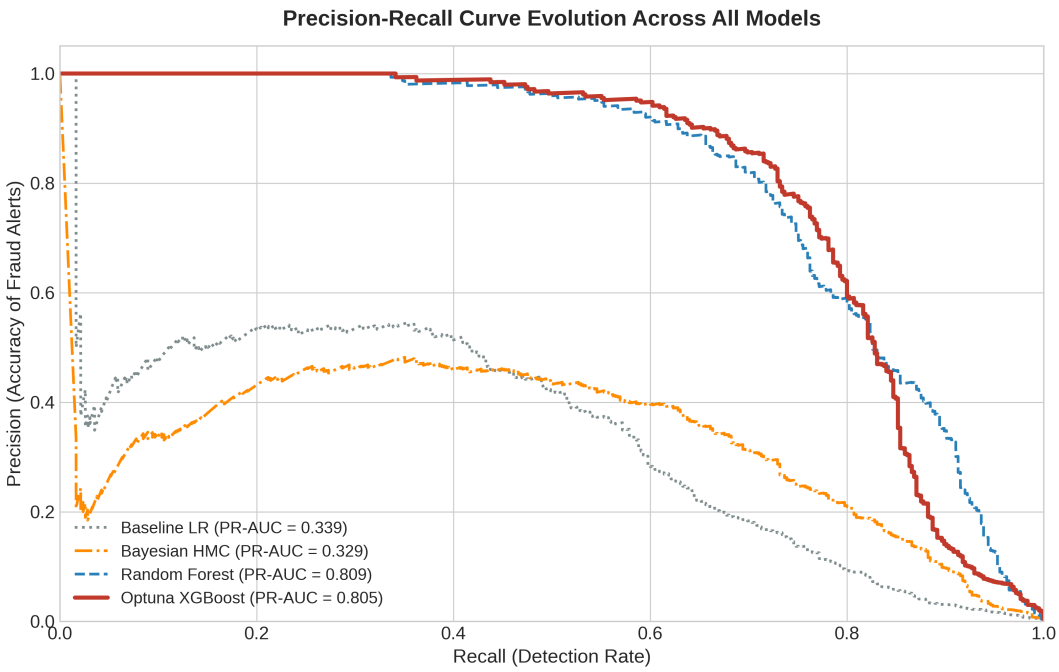

Figure 5: Precision-Recall Curve Evolution. The PR curve demonstrates that XGBoost architectures significantly outperform classical Logistic Regression. Notably, within the critical recall range of 0.5 to 0.8, the Optuna-optimized XGBoost maintains a superior precision-recall balance compared to the standard XGBoost. Furthermore, the Bayesian HMC model exhibits a clear improvement over the classical baseline LR.

As illustrated in Figure 5, tree-based ensembles—particularly the XGBoost architectures—significantly outperform the classical Logistic Regression baseline across all threshold levels. When balancing the trade-off between Precision and Recall, the hyperparameter-optimized XGBoost (Optuna XGBoost) demonstrates a distinct performance advantage over the standard XGBoost, especially within the critical Recall range of $[0.5, 0.8]$, where

the standard model begins to experience a steeper drop in precision. Furthermore, the transition to a probabilistic approach yields measurable benefits, as evidenced by the Cost-Sensitive Bayesian HMC model consistently maintaining higher precision than the classical Baseline LR.

4.4 Feature Importance and Business Interpretability

To interpret the decision-making logic of the optimal Optuna XGBoost model and ensure it aligns with real-world financial risk, we extracted the model's feature importance based on Information Gain. As visually depicted in Figure 6, the model is not acting as an opaque "black box"; rather, its learned rules strongly resonate with human intuition regarding fraud dynamics.

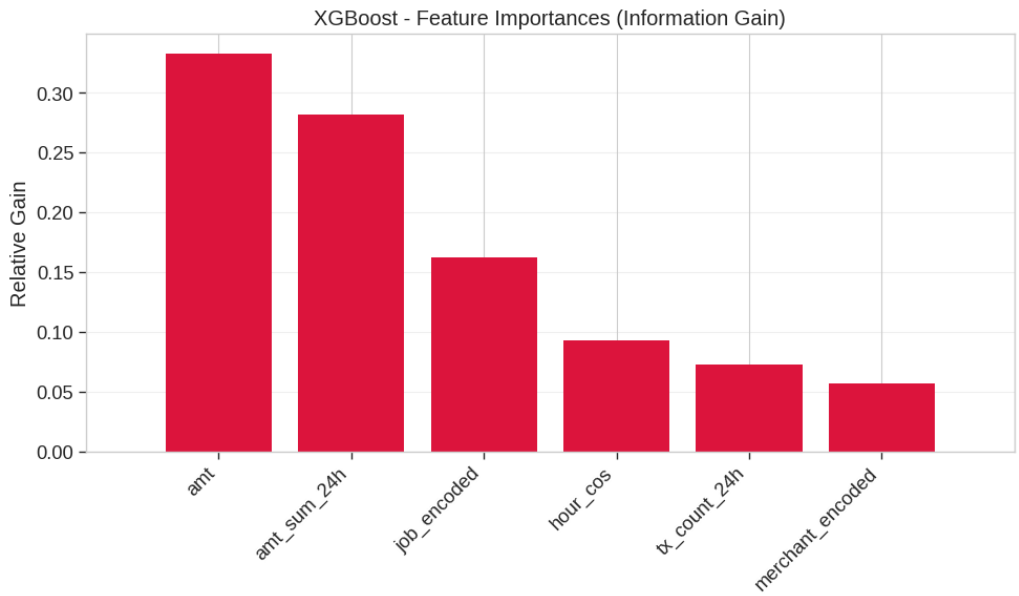

Figure 6: XGBoost Feature Importance (Information Gain). The model is overwhelmingly driven by monetary magnitude (`amt`) and rapid spending velocity (`amt_sum_24h`), directly validating the core assumptions of our cost-sensitive framework.

The analysis reveals that the transaction amount (`amt`) is the single most critical determinant, directly validating our cost-sensitive weight formulation. By explicitly penalizing high-value errors, the algorithm successfully learned to prioritize the magnitude of the transaction over trivial behavioral deviations.

Interestingly, the second most important feature is the rolling 24-hour expenditure (`amt_sum_24h`), which significantly outranks the sheer number of transactions (`tx_count_24h`). This provides a profound business insight: the model recognizes that fraudsters are highly motivated to drain a compromised credit limit as quickly as possible (spending velocity), rather than simply making numerous small purchases.

Furthermore, demographic and temporal risk indicators play a vital supplementary role. The high ranking of the Bayesian target-encoded job feature (`job_encoded`) suggests that specific socio-economic profiles are disproportionately targeted by cybercriminals.

Meanwhile, the prominent placement of the cyclical time encoding (`hour_cos`) mathematically confirms our initial Exploratory Data Analysis (EDA) findings from Section 2, validating that nocturnal operations—such as late-night transaction spikes—are critical signatures of fraudulent behavior. By translating these complex tree splits into actionable business intelligence, the model provides highly interpretable rules that can be seamlessly integrated into a bank's frontline risk-management engine.

5 Conclusion

This research presented a comprehensive evolutionary framework for credit card fraud detection, specifically addressing the dual challenges of extreme class imbalance and asymmetric misclassification costs. By transitioning from a classical linear baseline to a cost-sensitive Bayesian optimized XGBoost ensemble, we demonstrated that the integration of business-aligned penalties directly into the model's objective function yields superior results in protecting financial assets.

Our experimental findings highlight two critical insights for the FinTech industry. First, the application of a log-smooth cost-sensitive weight, defined as $5.0 + \ln(1 + \text{amt})$, effectively forces the model to prioritize high-value fraudulent transactions without sacrificing the detection of smaller "test" frauds. Second, we proved that a rigid adherence to the default 0.5 decision threshold is suboptimal for risk management. By strategically lowering the threshold to 0.30 for advanced tree ensembles, we achieved a remarkable 95.90% Protected Ratio, ensuring that nearly all fraudulent capital is intercepted while maintaining an operational precision of approximately 50%, which is well within the capacity of manual investigation teams.

5.1 Limitations and Future Work

Despite the robust performance of the current system, certain limitations remain. The current cost-sensitive weights are static and primarily based on the transaction amount; future iterations could incorporate dynamic cost functions that account for real-time merchant risk scores or customer lifetime value. Furthermore, while XGBoost excels at capturing tabular feature interactions, it does not inherently model the relational topology between entities.

Future research will explore the integration of *Graph Neural Networks (GNNs)* to capture complex money-laundering patterns and social-network-based fraud rings. Additionally, we aim to deploy this framework within a *Stream Processing* environment to evaluate its inference latency and adaptive learning capabilities in the face of rapidly evolving adversarial fraud tactics.

References

[1] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic Minority Over-sampling Technique," *Journal of Artificial Intelligence Research*, vol. 16, pp. 321-357, 2002.

[2] C. Elkan, "The foundations of cost-sensitive learning," in *Proceedings of the 17th International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 973-978, 2001.

[3] T. Chen and C. Guestrin, "XGBoost: A Scalable Tree Boosting System," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 785-794, 2016.

[4] T. Akiba, S. Sano, T. Yanase, T. Ohta, and M. Koyama, "Optuna: A Next-generation Hyperparameter Optimization Framework," in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 2623-2631, 2019.

[5] A. Scrivano, "Fraud Detection Pipeline Using Machine Learning: Methods, Applications, and Future Directions," Politecnico di Milano, DEIB, Italy, Preprint.

[6] S. Sundaravadivel, "Automating Fraud Detection: Machine Learning for Credit Card Security," ISMA University of Applied Sciences, Latvia, Preprint.